

# Github Repo

<https://github.com/Deep-Learning-Group-B/Case-Study-2>

```
In [ ]: import os
import json

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from tqdm import tqdm_notebook

import tensorflow as tf
from tensorflow.keras.layers import Dense, Reshape, Flatten, Dropout, LeakyReLU
from tensorflow.keras.layers import Conv2D, Conv2DTranspose
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.preprocessing.image import ImageDataGenerator, img_to_array

import warnings
warnings.filterwarnings('ignore')
```

## Load processed data

```
In [ ]: # Loading images from file
images = np.load('/content/drive/MyDrive/sem3_DL/images.npy')

# Loading labels from file
labels = np.load('/content/drive/MyDrive/sem3_DL/labels.npy')

print(f"Images loaded: {images.shape}")
print(f"Labels loaded: {labels.shape}")
```

Images loaded: (8390, 224, 224, 3)  
Labels loaded: (8390,)

```
In [ ]: labels_unique = np.unique(labels)
print(f"Unique labels: {labels_unique}")

Unique labels: ['drink' 'food' 'inside' 'menu' 'outside']
```

```
In [ ]: # Separate images by class
selected_class = 'food' # Replace with the actual class label you want to use
class_indices = np.where(labels == selected_class)[0]
class_images = images[class_indices]

print(f"Total images for class '{selected_class}': {class_images.shape}")
```

Total images for class 'food': (1678, 224, 224, 3)

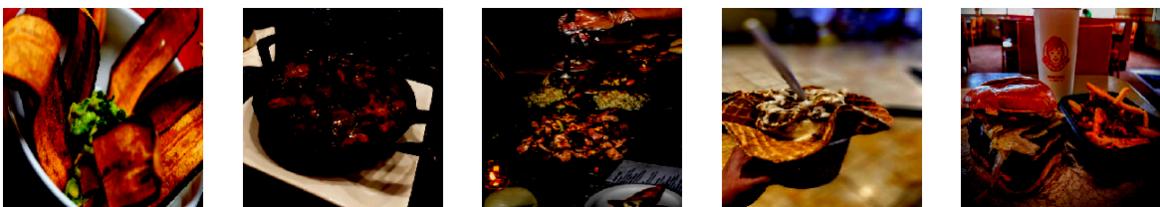
```
In [ ]: # Plot the images
num_images = 5
fig, axs = plt.subplots(1, num_images, figsize=(15, 15))
```

```

for i in range(num_images):
    axs[i].imshow(class_images[i, :, :, :]) # Display RGB images
    axs[i].axis('off')
plt.show()

```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).  
 WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).  
 WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).  
 WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).  
 WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).  
 WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



```

In [ ]: # Generator
def build_generator():
    model = Sequential()
    model.add(Dense(512 * 14 * 14, activation="relu", input_dim=100))
    model.add(Reshape((14, 14, 512)))
    model.add(Conv2DTranspose(256, kernel_size=4, strides=2, padding="same"))
    model.add(LeakyReLU(alpha=0.01))
    model.add(Conv2DTranspose(128, kernel_size=4, strides=2, padding="same"))
    model.add(LeakyReLU(alpha=0.01))
    model.add(Conv2DTranspose(64, kernel_size=4, strides=2, padding="same"))
    model.add(LeakyReLU(alpha=0.01))
    model.add(Conv2DTranspose(3, kernel_size=4, strides=2, activation="tanh", padding="same"))
    return model

# Discriminator
def build_discriminator():
    model = Sequential()
    model.add(Conv2D(64, kernel_size=4, strides=2, input_shape=(224, 224, 3), padding="same"))
    model.add(LeakyReLU(alpha=0.01))
    model.add(Dropout(0.3))
    model.add(Conv2D(128, kernel_size=4, strides=2, padding="same"))
    model.add(LeakyReLU(alpha=0.01))
    model.add(Dropout(0.3))
    model.add(Conv2D(256, kernel_size=4, strides=2, padding="same"))
    model.add(LeakyReLU(alpha=0.01))
    model.add(Dropout(0.3))
    model.add(Flatten())
    model.add(Dense(1, activation='sigmoid'))
    return model

```

```

In [ ]: # Compile the discriminator
discriminator = build_discriminator()
discriminator.compile(loss='binary_crossentropy', optimizer=Adam(), metrics=['accuracy'])

# Build and compile the combined model
generator = build_generator()
z = tf.keras.Input(shape=(100,))

```

```
img = generator(z)
discriminator.trainable = False
valid = discriminator(img)
combined = tf.keras.Model(z, valid)
combined.compile(loss='binary_crossentropy', optimizer=Adam())
```

In [ ]: `generator.summary()`

Model: "sequential\_1"

Layer (type)	Output Shape
dense_1 (Dense)	(None, 100352)
reshape (Reshape)	(None, 14, 14, 512)
conv2d_transpose (Conv2DTranspose)	(None, 28, 28, 256)
leaky_re_lu_3 (LeakyReLU)	(None, 28, 28, 256)
conv2d_transpose_1 (Conv2DTranspose)	(None, 56, 56, 128)
leaky_re_lu_4 (LeakyReLU)	(None, 56, 56, 128)
conv2d_transpose_2 (Conv2DTranspose)	(None, 112, 112, 64)
leaky_re_lu_5 (LeakyReLU)	(None, 112, 112, 64)
conv2d_transpose_3 (Conv2DTranspose)	(None, 224, 224, 3)

Total params: 12,891,587 (49.18 MB)

Trainable params: 12,891,587 (49.18 MB)

Non-trainable params: 0 (0.00 B)

In [ ]: `discriminator.summary()`

Model: "sequential"

Layer (type)	Output Shape
conv2d (Conv2D)	(None, 112, 112, 64)
leaky_re_lu (LeakyReLU)	(None, 112, 112, 64)
dropout (Dropout)	(None, 112, 112, 64)
conv2d_1 (Conv2D)	(None, 56, 56, 128)
leaky_re_lu_1 (LeakyReLU)	(None, 56, 56, 128)
dropout_1 (Dropout)	(None, 56, 56, 128)
conv2d_2 (Conv2D)	(None, 28, 28, 256)
leaky_re_lu_2 (LeakyReLU)	(None, 28, 28, 256)
dropout_2 (Dropout)	(None, 28, 28, 256)
flatten (Flatten)	(None, 200704)
dense (Dense)	(None, 1)

Total params: 859,585 (3.28 MB)

Trainable params: 0 (0.00 B)

Non-trainable params: 859,585 (3.28 MB)

In [ ]: `combined.summary()`

Model: "functional\_20"

Layer (type)	Output Shape
input_layer_2 (InputLayer)	(None, 100)
sequential_1 (Sequential)	(None, 224, 224, 3)
sequential (Sequential)	(None, 1)

Total params: 13,751,172 (52.46 MB)

Trainable params: 12,891,587 (49.18 MB)

Non-trainable params: 859,585 (3.28 MB)

```
# Function to generate and plot images
def generate_and_plot_images(generator, num_images=5):
    # Generate noise and use the generator to create images
    noise = np.random.normal(0, 1, (num_images, 100))
    generated_images = generator.predict(noise)

    # Rescale images 0 - 1
    generated_images = 0.5 * generated_images + 0.5
```

```
# Plot the images
fig, axs = plt.subplots(1, num_images, figsize=(15, 15))
for i in range(num_images):
    axs[i].imshow(generated_images[i, :, :, :]) # Display RGB images
    axs[i].axis('off')
plt.show()
```

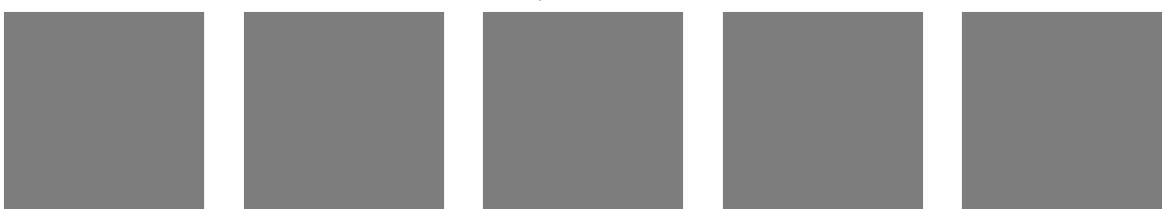
```
In [ ]: # Training parameters
epochs = 100
batch_size = 64
save_interval = 50

# Training Loop
for epoch in range(epochs):
    # Train Discriminator
    idx = np.random.randint(0, class_images.shape[0], batch_size)
    real_imgs = class_images[idx]
    noise = np.random.normal(0, 1, (batch_size, 100))
    fake_imgs = generator.predict(noise)
    d_loss_real = discriminator.train_on_batch(real_imgs, np.ones((batch_size, 1)))
    d_loss_fake = discriminator.train_on_batch(fake_imgs, np.zeros((batch_size, 1)))
    d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)

    # Train Generator
    noise = np.random.normal(0, 1, (batch_size, 100))
    g_loss = combined.train_on_batch(noise, np.ones((batch_size, 1)))

    # Print the progress
    if epoch % save_interval == 0:
        print(f"Epoch {epoch} [D loss: {d_loss[0]}] [G loss: {g_loss}]")
        generate_and_plot_images(generator, num_images=5)
```

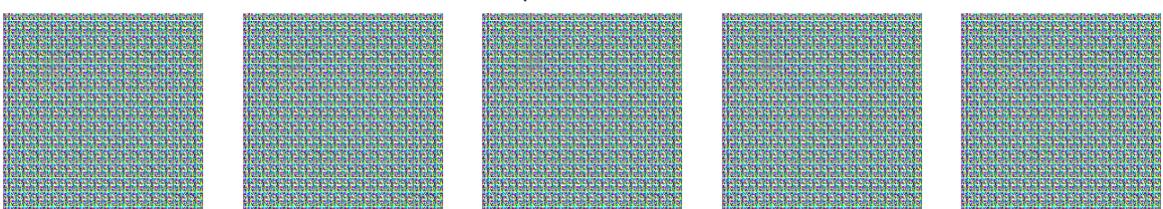
2/2 ————— 3s 4ms/step  
0 [D loss: 0.6547226309776306] [G loss: [array(0.6675807, dtype=float32), array(0.6675807, dtype=float32), array(0.4765625, dtype=float32)]]  
1/1 ————— 1s 892ms/step



2/2 ————— 0s 3ms/step

WARNING:tensorflow:5 out of the last 5 calls to <function TensorFlowTrainer.make\_train\_function.<locals>.one\_step\_on\_iterator at 0x7def680ed7e0> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce\_retracing=True option that can avoid unnecessary retracing. For (3), please refer to [https://www.tensorflow.org/guide/function#controlling\\_retracing](https://www.tensorflow.org/guide/function#controlling_retracing) and [https://www.tensorflow.org/api\\_docs/python/tf/function](https://www.tensorflow.org/api_docs/python/tf/function) for more details.

WARNING:tensorflow:6 out of the last 6 calls to <function TensorFlowTrainer.make\_train\_function.<locals>.one\_step\_on\_iterator at 0x7def680efc70> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce\_retracing=True option that can avoid unnecessary retracing. For (3), please refer to [https://www.tensorflow.org/guide/function#controlling\\_retracing](https://www.tensorflow.org/guide/function#controlling_retracing) and [https://www.tensorflow.org/api\\_docs/python/tf/function](https://www.tensorflow.org/api_docs/python/tf/function) for more details.



```
In [ ]: generator.save('/content/drive/MyDrive/sem3_DL/deployment/DCGAN_generator_food.h5')
discriminator.save('/content/drive/MyDrive/sem3_DL/deployment/DCGAN_discriminator.h5')
combined.save('/content/drive/MyDrive/sem3_DL/deployment/DCGAN_combined_food.h5')
```

```
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `ker  
as.saving.save_model(model)`. This file format is considered legacy. We recommend  
using instead the native Keras format, e.g. `model.save('my_model.keras')` or `ke  
ras.saving.save_model(model, 'my_model.keras')`.  
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `ker  
as.saving.save_model(model)`. This file format is considered legacy. We recommend  
using instead the native Keras format, e.g. `model.save('my_model.keras')` or `ke  
ras.saving.save_model(model, 'my_model.keras')`.  
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `ker  
as.saving.save_model(model)`. This file format is considered legacy. We recommend  
using instead the native Keras format, e.g. `model.save('my_model.keras')` or `ke  
ras.saving.save_model(model, 'my_model.keras')`.
```

In [ ]: